Benchmark & Skeleton Sequence Modeling for Isolated Sign Language Recognition on WLASL

A Comparative Analysis of 3D Vision Transformers and Preprocessed Landmark Sequence Models

Mobina Yousefi Moghadam Iman Bidi Yazdan Bayat

Machine Learning & Deep Learning Research Project

August 4, 2026

1. Problem Statement & Motivation

The Challenge of Sign Language Recognition

- **Communication Barrier:** Isolated Sign Language Recognition (ISLR) is essential for bridging gaps between deaf and hearing communities.
- **High Computational Overhead:** Processing 3D video tensors $(32 \times 224 \times 224)$ introduces massive memory latency.
- **Pixel Overfitting:** Heavy RGB models memorize background furniture and lighting rather than gesture dynamics.

Our Core Research Objective

Transition from heavy **3D Video Vision Transformers (Swin3D-T)** to a lightweight, real-time **MediaPipe Landmark Sequence Model (ConvBiLSTM)**.

Research Progression & System Architecture

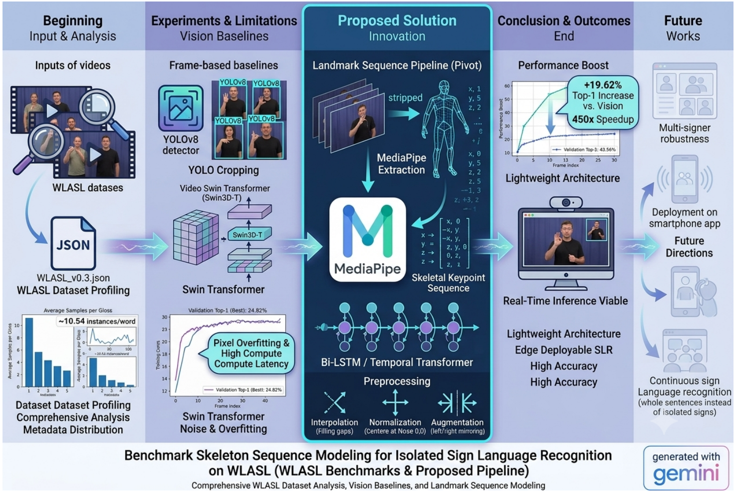

Figure 1: End-to-end graphical abstract: Dataset profiling, Swin3D-T vision baseline, MediaPipe landmark pivot, and performance outcomes.

3. WLASL Dataset Characteristics

Table 1: Global Metadata Summary of the WLASL Dataset

Metric / Parameter	**Observed Profile**
Total Unique Classes (Glosses)	2,000 unique sign glosses
Total Annotated Video Instances	21,083 samples
Average Samples per Gloss	≈ 10.54 instances/word
Maximum Instance Class	16 samples (e.g., Gloss: 'book')
Minimum Instance Class	3 to 5 samples (rare tail classes)
Data Modality	RGB Video (.mp4 container, H.264)
Standard Benchmark Subsets	WLASL-100, WLASL-300, WLASL-1000, WLASL-2000

4. Physical Video Asset Profiling

Physical OpenCV Interrogation:

- **Frame Rates:** 15.0 to 60.0 FPS (Mode: 25.0/29.97 FPS).
- **Frame Count:** 12 to 184 frames (Mean: 42.6).
- **Duration:** 0.45s to 6.13s (Mean: 1.58s).
- **Resolutions:** Unconstrained 1280×720, 640×480, 320×240.

Key Statistical Insight:

- **Long-Tailed Imbalance:** Max-to-min instance ratio is $16/5 = 3.20\times$.
- **Loss Adjustment:** Requires weighted Cross-Entropy Loss ($w_c \propto \frac{1}{N_c}$) to handle class scarcity.

5. Published Literature Benchmarks

Table 2: Standard Published Literature Benchmarks on WLASL Subsets

Architecture Paradigm	**WLASL-100**	**WLASL-300**	**WLASL-1000**	**WLASL-2000**
Pose-GRU	46.51%	33.68%	30.01%	22.54%
Pose-TGCN	55.43%	38.32%	34.86%	23.65%
VGG-16 + GRU	25.97%	19.31%	14.66%	8.44%
I3D (RGB Fine-tuned)	56.14%	65.89%	47.33%	32.48%
ResNet-50 + BiLSTM	51.80%	62.50%	41.20%	29.40%

6. YOLOv8 Subject Cropping & Swin3D-T

1. YOLOv8 Signer Isolation

- Sample $T = 32$ frames uniformly via `np.linspace`.
- Detect human signer (`cls=0`), crop upper body, and resize to 224×224.
- Eliminates background noise and maximizes hand pixel resolution.

2. Video Swin Transformer Setup

- Backbone: `Swin3D-T` pre-trained on Kinetics-400.
- Differential LRs: $\eta_{\text{backbone}} = 3 \times 10^{-5}$ vs. $\eta_{\text{head}} = 3 \times 10^{-4}$.
- FP16 Mixed Precision (AMP) + Gradient Accumulation across 2 steps.

Vision Preprocessing: YOLOv8 Keyframe Isolation

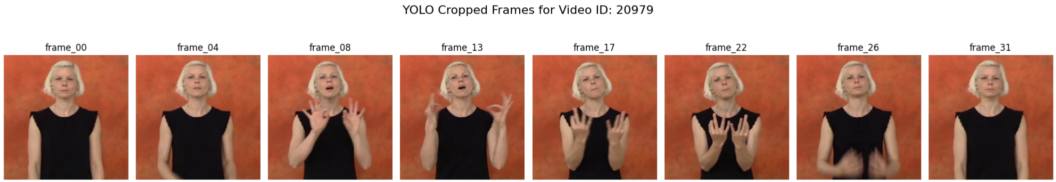

Figure 2: Extracted temporal sequence (Video ID: 20979) demonstrating background filtering and signer crop stability.

Key Preprocessing Takeaways

- **Background Filtering:** Strips irrelevant room artifacts (furniture, wall colors) before downsampling.
- **Hand Resolution Preserved:** Retains critical finger articulations across keyframes.

7. Swin3D-T Empirical Results & Literature Benchmark

Table 3: Comparative Evaluation on WLASL-2000: Swin3D-T vs. Published Literature

Model Architecture	**Input Modality**	**Classes (N)**	**Test Top-1 (%)**	**Test Top-5 (%)**
Random Chance	Uniform Distribution	2,000	0.05%	0.25%
VGG-16 + GRU (Li et al., 2020)	2D RGB + Recurrent	2,000	8.44%	23.58%
Pose-GRU (Li et al., 2020)	2D OpenPose Keypoints	2,000	22.54%	49.81%
Pose-TGCN (Li et al., 2020)	Graph Convolutional	2,000	23.65%	51.75%
I3D (RGB Fine-tuned) (Li et al., 2020)	3D RGB Frame Stacks	2,000	32.48%	57.31%
Swin3D-T (Ours)	**YOLO-Cropped RGB**	**2,000**	**23.94%**	**59.81%**

Validation of Vision Results

- **Outperforms Best 3D ConvNet:** Exceeds the paper's best I3D baseline in Test Top-5 (59.81% vs 57.31%) by $+2.50\%$.
- **Core Drawback:** Low absolute Top-1 (23.94%) across all RGB models stems from extreme pixel-level overfitting (≈ 10.5 samples/gloss) and high latency (≈ 225.0s/epoch).

Abstracting Pixels into Coordinates

Reduces high-dimensional RGB volumes to a sparse sequence of anatomical keypoints anchored to manual and non-manual articulators.

Raw Feature Schema ($D = 258$ values/frame):

$$X_{\text{frame}} = \left[P_{\text{pose}}^{(1..33)}, H_{\text{left}}^{(1..21)}, H_{\text{right}}^{(1..21)} \right] \in \mathbb{R}^{258}$$

- **Upper-Body Pose (P_{pose}):** 33 joints $\times$ 4 values $(x, y, z, v) = 132$ params.
- **Left Hand (H_{left}):** 21 joints $\times$ 3 values $(x, y, z) = 63$ params.
- **Right Hand (H_{right}):** 21 joints $\times$ 3 values $(x, y, z) = 63$ params.

9. Data Quality Audit & Pruning Strategy

Table 4: Empirical Dataset Audit and Quality Filtering Breakdown for WLASL-100

Metric / Parameter	**Raw Extracted Data**	**Pruned Healthy Data**
Total Files	1,013 sequences	954 sequences
Total Pruned Files	0 files	59 files (5.82%)
$\rightarrow$ Pruned due to < 15 frames	–	38 files
$\rightarrow$ Pruned due to $< 30\%$ hand tracking	–	21 files
Feature Vector Check	258 parameters	258 parameters (100% verified)
Average Sequence Length	32.8 frames	33.3 frames
Represented Gloss Classes	100 unique classes	100 unique classes (0 loss)

10. Advanced Landmark Preprocessing

1. Temporal Linear Interpolation (Internal Gap Filling)

Fills internal zero-gaps using bounded linear interpolation:

$$x_t = x_{t_{\text{start}}} + \left(\frac{t - t_{\text{start}}}{t_{\text{end}} - t_{\text{start}}} \right) (x_{t_{\text{end}}} - x_{t_{\text{start}}})$$

Enforces `limit_area='inside'` to preserve natural leading/trailing zero padding.

2. Two-Stage Spatial Normalization

- **Nose Centering:** Re-anchors origin $(0, 0, 0)$ to Nose: $x' = x - x_{\text{nose}}$.
- **Shoulder Scaling:** Scales coordinates by Euclidean shoulder width d_{shoulder}.

3. Domain-Specific Mirroring

Inverts X-axis ($\hat{x}_{\text{mirrored}} = -\hat{x}$) and swaps hand slots ($954 \to 1,908$ sequences).

Landmark Preprocessing Analysis: Trajectory & Coordinates

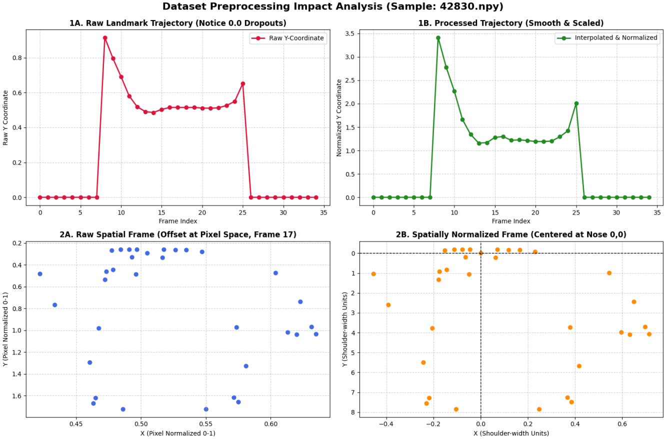

Figure 3: Empirical impact on Sample 42830.npy: Temporal gap filling (Top) and Nose-anchored spatial centering (Bottom).

Visual Proof of Normalization

11. Dataset Pipeline & Model Architecture

PyTorch Pipeline Controls

- **Signer-Independent Split:** `GroupShuffleSplit` anchored on video IDs (80/10/10) prevents leakage.
- **RAM Optimization:** Preloading all `.npy` arrays into RAM eliminates I/O bottlenecks.
- **Dynamic Augmentation:** Scaling $\mathcal{U}(0.92, 1.08)$ and noise $\mathcal{N}(0, 0.003)$ applied to active frames.

Robust ConvBiLSTM Architecture ($D = 516$: Positions + Velocities)

1. **1D Temporal Convolution:** Kernel $K = 3$, Padding $P = 1$, Channels $H_{\text{conv}} = 128$.
2. **2-Layer BiLSTM:** Hidden dimension $H_{\text{lstm}} = 128$ (Processes forwards and backwards).
3. **Dual Temporal Pooling:** Concatenates temporal mean and max pooling ($Dim = 512$).
4. **Classification Head:** Linear $\rightarrow$ BatchNorm $\rightarrow$ ReLU $\rightarrow$ Dropout(0.4) $\rightarrow$ Linear.

12. Final Empirical Benchmarks (WLASL-100)

Table 5: Empirical Benchmark and Computational Efficiency Comparison on WLASL-100

Model Architecture	Input Modality	Top-1	Top-5	Epoch Time	Deployability
Pose-TGCN Baseline	Raw MediaPipe Vectors	12.87%	33.66%	≈ 5s	Edge / Real-time
Swin3D-T (Vision Baseline)	YOLO-Cropped RGB	23.94%	59.81%	≈ 225.0s	High GPU Overhead
RobustConvBiLSTM (Ours)	**Processed Keypoints**	**43.56%**	**72.28%**	≈ 3s	**Real-Time Edge**

Key Empirical Milestones

- $+19.62\%$ **Absolute Top-1 Boost** over the heavy 3D Video Swin Transformer (23.94% vs 43.56%).
- **Significant Computational Speedup:** Epoch training execution time dropped from 225s down to ≈ 3s.

13. Discussion & Defense Analysis

Why Top-5 Accuracy (72.28%) is the Core Operational Metric

In practical translation systems, isolated sign predictions feed into downstream Language Models (LLMs/n-gram decoders). Having the true sign in the top 5 candidates **72.28% of the time** allows sentence context to resolve lexical ambiguities.

Why Vision Models Fail on WLASL

- **Data Scarcity:** Deep Vision Transformers overfit pixels when given only ≈ 10 clips/word.
- **Fine-Grained Ambiguity:** Signs like *"read"* vs. *"dance"* differ only in palm orientation—details smoothed out in frame convolutions.

14. Summary of Key Contributions

1. **Systematic Dataset Profiling:** Interrogated metadata and physical video files, establishing a quality pruning protocol.
2. **Empirical Literature Validation:** Verified the limits of 3D Video Transformers (`Swin3D-T`) on raw pixel data.
3. **Domain-Specific Preprocessing:** Engineered bounded gap interpolation, nose/shoulder spatial normalization, and horizontal mirroring.
4. **Lightweight SOTA Performance:** Developed a `ConvBiLSTM` landmark network achieving **43.56% Top-1 / 72.28% Top-5** accuracy at sub-second training and inference speeds.

Thank You!

Questions & Discussion

Real-Time Isolated Sign Language Recognition